

Problem Set 2 (PHYS641) - Matias Castro Tapia

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
```

Problem 1

I defined the `qrpo` routine to do the least squares solution for fit parameters (for polynomial fitting) using the QR decomposition of A . Then, if $A = QR$ and we have $A^T N^{-1} A m = A^T N^{-1} d$:

$$\begin{aligned} (QR)^T N^{-1} (QR) m &= (QR)^T N^{-1} d \\ (R^T Q^T N^{-1} QR) m &= R^T Q^T N^{-1} d \\ m &= (R^T Q^T N^{-1} QR)^{-1} R^T Q^T N^{-1} d \\ m &= R^{-1} (Q^T N^{-1} Q)^{-1} R^{T,-1} R^T Q^T N^{-1} d \\ m &= R^{-1} (Q^T N^{-1} Q)^{-1} Q^T N^{-1} d \end{aligned}$$

For the specific case where $N = I$ we will have:

$$m = R^{-1} (Q^T Q)^{-1} Q^T N^{-1} d = R^{-1} Q^T N^{-1} d$$

since $Q^T Q = I$. My routine considers a general N anyway.

The routine receives x points and data $y = d$, the polynomial order o and, a noise matrix defined as the input $s = N$. The routine prints the model parameters and returns the data prediction, i.e., Am .

In [2]:

```
def qrpo(x,y,o,s):
    n=len(y)
    mat=np.zeros([n,o+1])
    for i in range(o+1):
        mat[:,i]=x**i
    q,r=np.linalg.qr(mat)

    Ninv=np.linalg.inv(s)
    mod=np.linalg.inv(r).T@np.linalg.inv(q.T@Ninv@q).T@Ninv@y
    print(mod)
    pred=mat@mod
    return pred
```

I also defined the routine `atpo` for solving the least squares without the QR decomposition. Then, the solution is just $m = (A^T N^{-1} A)^{-1} A^T N^{-1} d$

In [3]:

```
def atpo(x,y,o,s):
    n=len(y)
    mat=np.zeros([n,o+1])
    for i in range(o+1):
        mat[:,i]=x**i

    Ninv=np.linalg.inv(s)
    mod=np.linalg.inv(mat.T@Ninv@mat)@mat.T@Ninv@y
    print(mod)
    pred=mat@mod
    return pred
```

Let's do a fit for $f(x) = \sin x \cos x$.

In [4]:

```
x=np.linspace(-3,3,301)
```

In [5]:

```
y1=np.sin(x)*np.cos(x)
```

I defined a noise matrix as an identity matrix.

In [6]:

```
n=len(x)
N=np.zeros([n,n])
for i in range(n):
    N[i][i]=1
```

A 22th order polynomial fit using both methods.

In [7]:

```
y1predq=qrpo(x,y1,22,N)
```

```
[ 4.01940850e-16  1.00000000e+00 -6.18342183e-15 -6.66666667e-01
 2.31498848e-14  1.33333333e-01 -4.33091063e-14 -1.26984121e-02
 3.95100619e-14  7.05466894e-04 -2.18280255e-14 -2.56531355e-05
 7.54691448e-15  6.57711246e-07 -1.64717415e-15 -1.25158437e-08
 2.31504269e-16  1.82533983e-10 -2.02201588e-17 -2.01367224e-12
 1.00572258e-18  1.36829224e-14 -2.14070207e-20]
```


In [8]:

```
y1preda=atpo(x,y1,22,N)
```

```
[-2.77041326e-04  1.00204414e+00  1.70631879e-04 -6.67263895e-01
 -4.26173355e-05  1.33472485e-01  9.13443838e-06 -1.27162315e-02
 -4.92405347e-06  7.01804660e-04  2.41763747e-06 -2.36882014e-05
 -7.51280407e-07  2.22403810e-07  1.50679239e-07  4.85946147e-08
 -1.95582349e-08 -5.20424665e-09  1.58368688e-09  2.65643865e-10
 -7.26439171e-11 -5.69363998e-12  1.44354489e-12]
```

In [9]:

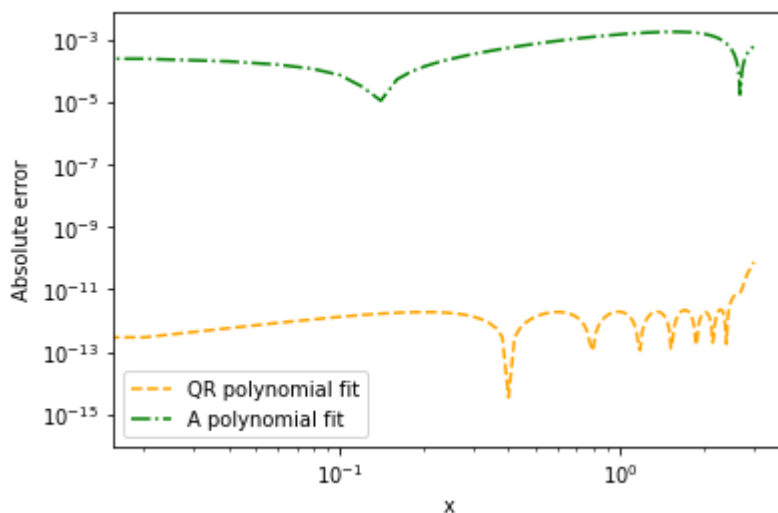
```
plt.plot(x,y1,label='$f(x)=\sin{(x)} \cos{(x)}$')
plt.plot(x,y1predq,'--',label='QR polynomial fit')
plt.plot(x,y1preda,'-.',label='A polynomial fit')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.show()
```



It seems like both fits are very close to the function, so let's the absolute error:

In [10]:

```
plt.loglog(x,np.abs(y1predq-y1),'--',label='QR polynomial fit',color='orange')
plt.loglog(x,np.abs(y1preda-y1),'-.',label='A polynomial fit',color='green')
plt.xlabel('x')
plt.ylabel('Absolute error')
plt.legend()
plt.show()
```



Then, the absolute error is much smaller when using the QR decomposition.

Problem 2

I defined the routine `Tn_fit` to do a Chebyshev polynomial fit for a set of points x , data y , and order od . I used the Chebyshev polynomial definition $T_n = \cos(n \arccos(x))$ for generating the n^{th} order polynomial. Then, the i^{th} row of the matrix A for the model must be $(T_0(x_i) \ T_1(x_i) \dots T_{od}(x_i))$ with x_i the i^{th} point. To the least squares fit for the coefficients I just used the classic linear solution $m = (A^T N^{-1} A)^{-1} A^T N^{-1} d = (A^T A)^{-1} A^T d$ (ignoring the noise matrix). The routine returns the prediction Am and the set of coefficients m .

In [11]:

```
def Tn_fit(x,y,od):
    n=len(x)
    mat=np.zeros([n,od+1])
    for i in range(od+1):
        mat[:,i]=np.cos(i*np.arccos(x))

    mod=np.linalg.inv(mat.T@mat)@mat.T@y

    pred=mat@mod
    return pred,mod
```

I also defined the routine `Tn_eval` to eval an array of x points to a model coefficients as:

$$y = m_0 T_0(x) + m_1 T_1(x) + \dots + m_{od} T_{od}(x)$$

where od =length of m .

In [12]:

```
def Tn_eval(x,m):
    ev=[]
    for i in range(len(m)):
        ev.append(m[i]*np.cos(i*np.arccos(x)))
    return np.array(sum(ev))
```

Let's take the `xx` array for the fit to e^x

In [13]:

```
xx=np.linspace(-1,1,1000)
```


Just checking that $T_n = \cos(n \arccos(x))$ works and that they are polynomials.

In [14]:

```

for i in range(7):
    plt.plot(xx,np.cos(i*np.arccos(xx)))
plt.xlabel('x')
plt.ylabel('$T_{n}(x)$')
#plt.legend()
plt.show()

```



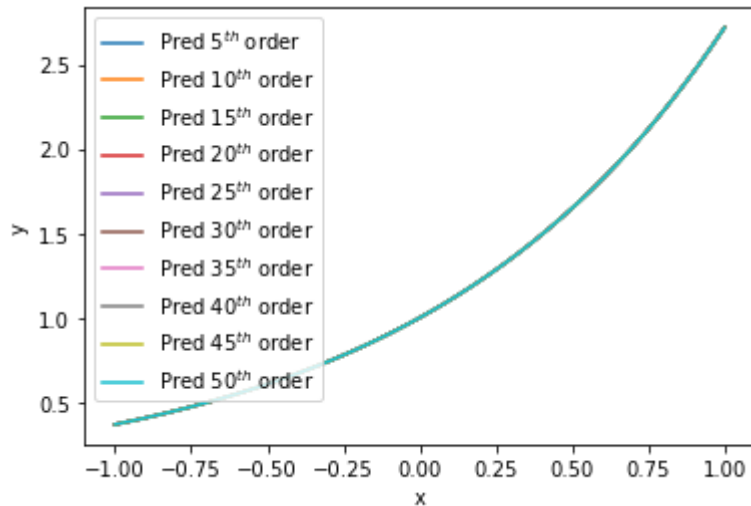
In [15]:

```
yy=np.exp(xx)
```

Let's check the stability of the fit.

In [16]:

```
for i in range(5,55,5):  
    predt,c=Tn_fit(xx,yy,i)  
    plt.plot(xx,predt,label='Pred '+str(i)+'$^{th}$ order')  
plt.xlabel('x')  
plt.ylabel('y')  
plt.legend()  
plt.show()
```



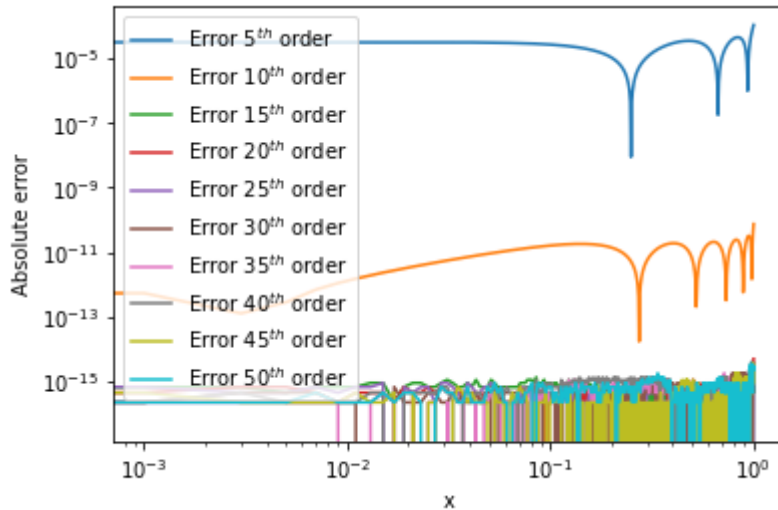
It seems like it remains very stable even for a 50th order fitting. So, we should look the absolute errors too.

In [17]:

```

for i in range(5,55,5):
    predt,c=Tn_fit(xx,yy,i)
    plt.loglog(xx,np.abs(yy-predt),label='Error '+str(i)+'$^{th}$ order')
plt.xlabel('x')
plt.ylabel('Absolute error')
plt.legend()
plt.show()

```



The error remains near 10^{-15} from about the 15th order. Let's do a 6th and a 30th order fit for part b).

In [18]:

```

tn6,c6=Tn_fit(xx,yy,6)
print(c6)

```

```

[1.26606584e+00 1.13031696e+00 2.71495266e-01 4.43354952e-02
 5.47415863e-03 5.41243732e-04 4.48739566e-05]

```

In [19]:

```

maxe6=np.max(np.abs(tn6-yy))
rms6=np.sqrt(sum(np.abs(tn6-yy)**2)/len(yy))

```

In [20]:

```

print('Max error for 6th order: '+str(maxe6))
print('RMS error for 6th order: '+str(rms6))

```

```

Max error for 6th order: 7.98480781183386e-06
RMS error for 6th order: 1.9852969888235997e-06

```

In [21]:

```
tn30,c30=Tn_fit(xx,yy,30)
print(c30)
```

```
[ 1.26606588e+00  1.13031821e+00  2.71495340e-01  4.43368498e-02
 5.47424044e-03  5.42926312e-04  4.49773230e-05  3.19843646e-06
 1.99212481e-07  1.10367713e-08  5.50589491e-10  2.49789842e-11
 1.03899528e-12  3.93574062e-14  1.32532874e-15 -1.73472348e-16
-3.46944695e-18 -1.52655666e-16 -4.16333634e-17 -2.49800181e-16
 1.17961196e-16 -7.63278329e-17 -3.46944695e-18  1.04083409e-16
 4.51028104e-17 -3.81639165e-17  6.59194921e-17 -9.02056208e-17
 1.24900090e-16 -1.42247325e-16  3.29597460e-17]
```

Now, I have to truncate the 30th order fit to 6th, and evaluate the truncated fit using Tn_eval

In [22]:

```
c30[7:]=0
print(c30)
```

```
[1.26606588e+00 1.13031821e+00 2.71495340e-01 4.43368498e-02
 5.47424044e-03 5.42926312e-04 4.49773230e-05 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00
 0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
```

In [23]:

```
tn30tr=Tn_eval(xx,c30)
```

In [24]:

```
maxe30t=np.max(np.abs(tn30tr-yy))
rms30t=np.sqrt(sum(np.abs(tn30tr-yy)**2)/len(yy))
```

In [25]:

```
print('Max error for 30th order truncated fit: '+str(maxe30t))
print('RMS error for 30th order truncated fit: '+str(rms30t))
```

```
Max error for 30th order truncated fit: 3.4092623653059206e-06
RMS error for 30th order truncated fit: 2.2588113160560264e-06
```

In [26]:

```
maxe6/maxe30t,rms6/rms30t
```

Out[26]:

```
(2.342092498685523, 0.8789122733323323)
```

If we look the old coefficients of the 30th order fit (c30 before truncating), we can note that the 8th coefficient is very similar to the Max error on the truncated 30th order fit, both are about 3×10^{-6} .

On the other hand, when comparing the max error for the 6^{th} order fit and the truncated 30^{th} order, we can see that it was reduced by a factor of ~ 2.3 when we truncated, while the RMS error is larger for about 14%.

Problem 3

We can start with the Cholesky decomposition of a noise matrix, then $N = LL^T$. If we consider the eigendecomposition of the matrix $N = V\Lambda V^T$ and remember that if we know $N = LL^T$ and the errors are correlated we can go from correlated data to uncorrelated as follows:

$$d_{uncorr} = L^{-1}d_{corr}$$

Thus, we can consider that $L = V\Lambda^{1/2}$ for positive eigenvalues (because N is a positive defined matrix) based on the two type of decomposition discussed and we can obtain correlated data from uncorrelated data using the eigenvalues/vectors:

$$d_{corr} = Ld_{uncorr} = V\Lambda^{1/2}d_{uncorr}$$

I defined the routine `correlated_data` to generate random correlated data from an input of a noise matrix of correlated errors. The routine generates a set of random values (consistent with the length of the noise matrix) and find the eigenvalues/vectors of the noise matrix using `numpy.linalg.eig`. Then, it generates correlated data using $d_{corr} = V\Lambda^{1/2}d_{uncorr}$ and returns it along with $\langle d_{corr}d_{corr}^T \rangle$.

$V = [\vec{v}_1 \ \vec{v}_2 \dots \vec{v}_n]$ for \vec{v}_i every eigenvector as a column and the element $\Lambda_{ii} = \lambda_i$ and 0 for non-diagonal elements for every eigenvalue λ_i .

In [27]:

```
def correlated_data(Ncorr):
    n=len(Ncorr)
    d=np.random.randn(len(Ncorr))

    l,vec=np.linalg.eig(Ncorr)

    lam_1=np.zeros([n,n])
    for i in range(n):
        lam_1[i][i]=l[i]

    dcorr=vec@np.sqrt(lam_1)@d

    return dcorr,np.transpose([dcorr])@[dcorr]
```

I created a matrix $N_{ij} = 1 + \delta_{ij}$ of dimensions 5×5

In [28]:

```
Ncorr=np.ones([5,5])
for i in range(len(Ncorr)):
    Ncorr[i][i]=2
```

In [29]:

Ncorr

Out[29]:

```
array([[2., 1., 1., 1., 1.],
       [1., 2., 1., 1., 1.],
       [1., 1., 2., 1., 1.],
       [1., 1., 1., 2., 1.],
       [1., 1., 1., 1., 2.]])
```

I used the routine 100 times to obtain the average of the matrix $\langle d_{corr} d_{corr}^T \rangle$ for these 100 iterations.

In [30]:

```
dc100=[]
ddt100=[]
for i in range(100):
    dc,ddt=correlated_data(Ncorr)
    dc100.append(dc)
    ddt100.append(ddt)
```

In [31]:

```
sum(ddt100)/100
```

Out[31]:

```
array([[2.72615432, 1.09263198, 1.60634134, 1.3019616 , 1.69560458],
       [1.09263198, 2.57453737, 1.11429005, 1.28721382, 0.95478667],
       [1.60634134, 1.11429005, 2.26958561, 1.4525458 , 1.45165529],
       [1.3019616 , 1.28721382, 1.4525458 , 1.76675411, 1.40441788],
       [1.69560458, 0.95478667, 1.45165529, 1.40441788, 2.12578387]])
```

Now for 10000 iterations:

In [32]:

```
dc10000=[]
ddt10000=[]
for i in range(10000):
    dc,ddt=correlated_data(Ncorr)
    dc10000.append(dc)
    ddt10000.append(ddt)
```

In [33]:

```
sum(ddt10000)/10000
```

Out[33]:

```
array([[2.06883928, 0.75486194, 1.08752121, 0.93277717, 1.18082445],
       [0.75486194, 2.37963373, 0.93391176, 1.33569191, 0.68050593],
       [1.08752121, 0.93391176, 1.91416198, 1.11806224, 1.01789104],
       [0.93277717, 1.33569191, 1.11806224, 1.5830557 , 1.0546415 ],
       [1.18082445, 0.68050593, 1.01789104, 1.0546415 , 2.11083216]])
```

And 100000:

In [34]:

```
dc100000=[]
ddt100000=[]
for i in range(100000):
    dc,ddt=correlated_data(Ncorr)
    dc100000.append(dc)
    ddt100000.append(ddt)
```

In [35]:

```
sum(ddt100000)/100000
```

Out[35]:

```
array([[2.08437291, 0.75298867, 1.09358325, 0.94062985, 1.1803205 ],
       [0.75298867, 2.36025524, 0.91026667, 1.32259239, 0.67918843],
       [1.09358325, 0.91026667, 1.90795448, 1.11370475, 1.0218473 ],
       [0.94062985, 1.32259239, 1.11370475, 1.59607475, 1.06773362],
       [1.1803205 , 0.67918843, 1.0218473 , 1.06773362, 2.10052975]])
```

We can note that the average for many realizations $\langle d_{corr} d_{corr}^T \rangle \sim N$

Problem 4

I defined the routine `corr_noise` to generate the correlated noise matrix $N_{ij} = a \exp(\frac{-(i-j)^2}{2\sigma^2}) + (1-a)\delta_{ij}$ with $a, s = \sigma$, and x as inputs. x is just used to now the length of the matrix.

In [36]:

```
def corr_noise(x,a,s):
    n=len(x)
    N=np.zeros([n,n])
    for i in range(n):
        for j in range(n):
            if i==j:
                N[i][j]=a*np.exp(-((i-j)**2)/(2*(s**2)))+(1-a)
            else:
                N[i][j]=a*np.exp(-((i-j)**2)/(2*(s**2)))
    return N
```

In [37]:

```
xs=np.linspace(0,999,1000)
```


I defined the Gaussian signal for the x points.

In [38]:

```
signal=np.exp(-((xs-xs[499])**2)/(2*(50**2)))
```

In [39]:

```
plt.plot(xs,signal)
plt.xlabel('x')
plt.ylabel('amp')
#plt.legend()
plt.show()
```



I computed the error bar for the fit amplitude for each pair of a and σ using the output of the routine `corr_noise`. The estimation of the error bar is obtained as $\sqrt{(A^T N^{-1} A)^{-1}}$, with A the template signal and N the matrix of correlated noise.

In [40]:

```
a=[0.1,0.5,0.9]
sigma=[5,50,500]
for i in a:
    for j in sigma:
        error=np.sqrt(1/(signal@np.linalg.inv(corr_noise(xs,i,j))@signal))
        print('For a='+str(i)+' and sigma='+str(j)+
              ' the error bar on the fit is:'+str(error))
```

```
For a=0.1 and sigma=5 the error bar on the fit is:0.15576335487646098
For a=0.1 and sigma=50 the error bar on the fit is:0.3377448956231588
For a=0.1 and sigma=500 the error bar on the fit is:0.12757498758982014
For a=0.5 and sigma=5 the error bar on the fit is:0.27599935080333926
For a=0.5 and sigma=50 the error bar on the fit is:0.7140544770562327
For a=0.5 and sigma=500 the error bar on the fit is:0.10066392140726646
For a=0.9 and sigma=5 the error bar on the fit is:0.3578952599636376
For a=0.9 and sigma=50 the error bar on the fit is:0.9499319468584407
For a=0.9 and sigma=500 the error bar on the fit is:0.04892549466815368
```

The best error was for the case of $a = 0.9$ and $\sigma = 500$ where the data must be correlated over large distances but almost perfectly correlated; the error bar was 0.0489. The worst case was $a = 0.9$ and $\sigma = 50$ where the error bar was 0.9499, in this case a and σ are very closed to the amplitude and σ_{src} values, then

probably when estimating $\sqrt{(A^T N^{-1} A)^{-1}}$, the coupling of the source signal and the errors could cause that the error bars will be of the order of the amplitude. The second worst was $a = 0.5$ and $\sigma = 50$ where the error bar was 0.714, that also has $\sigma = \sigma_{src}$.